

— **AVOD** —

Aggregate View Object Detection





ONE

Kitti Object Detection Dataset

TWO

AVOD Algorithm

THREE

AVOD code




PART 1

Kitti Object Detection Dataset

http://www.cvlibs.net/datasets/kitti/eval_object.php?obj_benchmark=3d

Introduction



data collecting platform:

- $2 \times$ PointGray Flea2 grayscale cameras
- $2 \times$ PointGray Flea2 color cameras
- $1 \times$ Velodyne HDL-64E rotating 3D laser scanner
- $1 \times$ OXTS RT3003 inertial and GPS navigation system
- $4 \times$ Edmund Optics lenses

tasks of interest :

- stereo, optical flow, visual odometry, 3D object detection and 3D tracking.

Fig. 1. **Recording Platform.** Our VW Passat station wagon is equipped with four video cameras (two color and two grayscale cameras), a rotating 3D laser scanner and a combined GPS/IMU inertial navigation system.

3D Object Detection Dataset

class:

- 'Van', 'Car', 'Truck', 'Pedestrian', 'Person (sitting)', 'Cyclist', 'Tram', 'Misc'

3D bounding box overlap :

- Car:70%
- pedestrians 、 cyclists:50%


Difficulty:

- **Easy:** Min. bounding box height: 40 Px, Max. occlusion level: Fully visible, Max. truncation: 15 %
- **Moderate:** Min. bounding box height: 25 Px, Max. occlusion level: Partly occluded, Max. truncation: 30 %
- **Hard:** Min. bounding box height: 25 Px, Max. occlusion level: Difficult to see, Max. truncation: 50 %




Type	truncated	occluded	alpha	bbox				dimensions			location			rotation_y	score
Truck	0.74	1	2.07	0	0	424.74	374	2.6	2.06	5.42	-3.17	1.77	5.46	1.57	
Car	0	0	-1.81	742.41	184.49	944.56	321.39	1.46	1.6	3.71	2.84	1.63	9.72	-1.54	
Van	0	0	-1.64	639.17	169.69	683.48	212.97	1.97	1.82	4.41	2.42	1.86	35.34	-1.57	
Car	0	0	-1.48	551.01	184.06	575.42	204.29	1.58	1.65	3.91	-3.85	2.52	59.59	-1.55	
DontCare	-1	-1	-10	579.35	178.15	633.56	201.11	-1	-1	-1	-1000	-1000	-1000	-10	
DontCare	-1	-1	-10	527.27	181.27	543.98	207.35	-1	-1	-1	-1000	-1000	-1000	-10	

Point Cloud



- **N*4 matrices:**
(x,y,z, reflectivity)

- **Project to image coordinate:**
 $x = P2 * R0_rect * Tr_velo_to_cam * y$



Evaluation



$$a_o = \frac{\text{area } B_p \cap B_{gt}}{\text{area } B_p \cup B_{gt}}$$

$$\text{Precision} = \frac{TP}{TP + FP}$$

$$\text{Recall} = \frac{TP}{TP + FN}$$

- **Average Precision (AP) :**

$$AP = \frac{1}{11} \sum_{r \in \{0, 0.1, \dots, 1\}} p_{interp}(r)$$

- **Average Orientation Similarity (AOS) :**

$$AOS = \frac{1}{11} \sum_{r \in \{0, 0.1, \dots, 1\}} \max_{\hat{r}: \hat{r} \geq r} s(\hat{r})$$

$$s(r) = \frac{1}{|D(r)|} \sum_{i \in D(r)} \frac{1 + \cos \Delta_{\theta}^{(i)}}{2} \delta_i$$




PART 2

AVOD Algorithm

- Difficulties of 3d object detection
- AVOD architecture
 - generating feature maps
 - RPN network
 - second stage detection network

Compare



CAT, DOG, DUCK

1. algorithm:

Faster R- CNN, R-FCN, YOLO、 SSD

2. input:

RGB images

3.output:

2D boxes、 classes、 confidence



1. algorithm:

3DOP, Deep3DBox, VoxelNet, MV3D

2. input:

RGB images、 LIDAR

3.output:

2D boxes、 3D boxes、 classes、 orientation、 confidence

Difficulties of 3D Object Detection

- Needs to capture depth and orientation from environment
- More complex and computationally expensive processing at later detection stages
- Hard to detect smaller objects such as pedestrian and cyclists
- Methods only uses LIDAR point clouds, stereo depth maps or RGBD sensor depth maps

AVOD Architecture

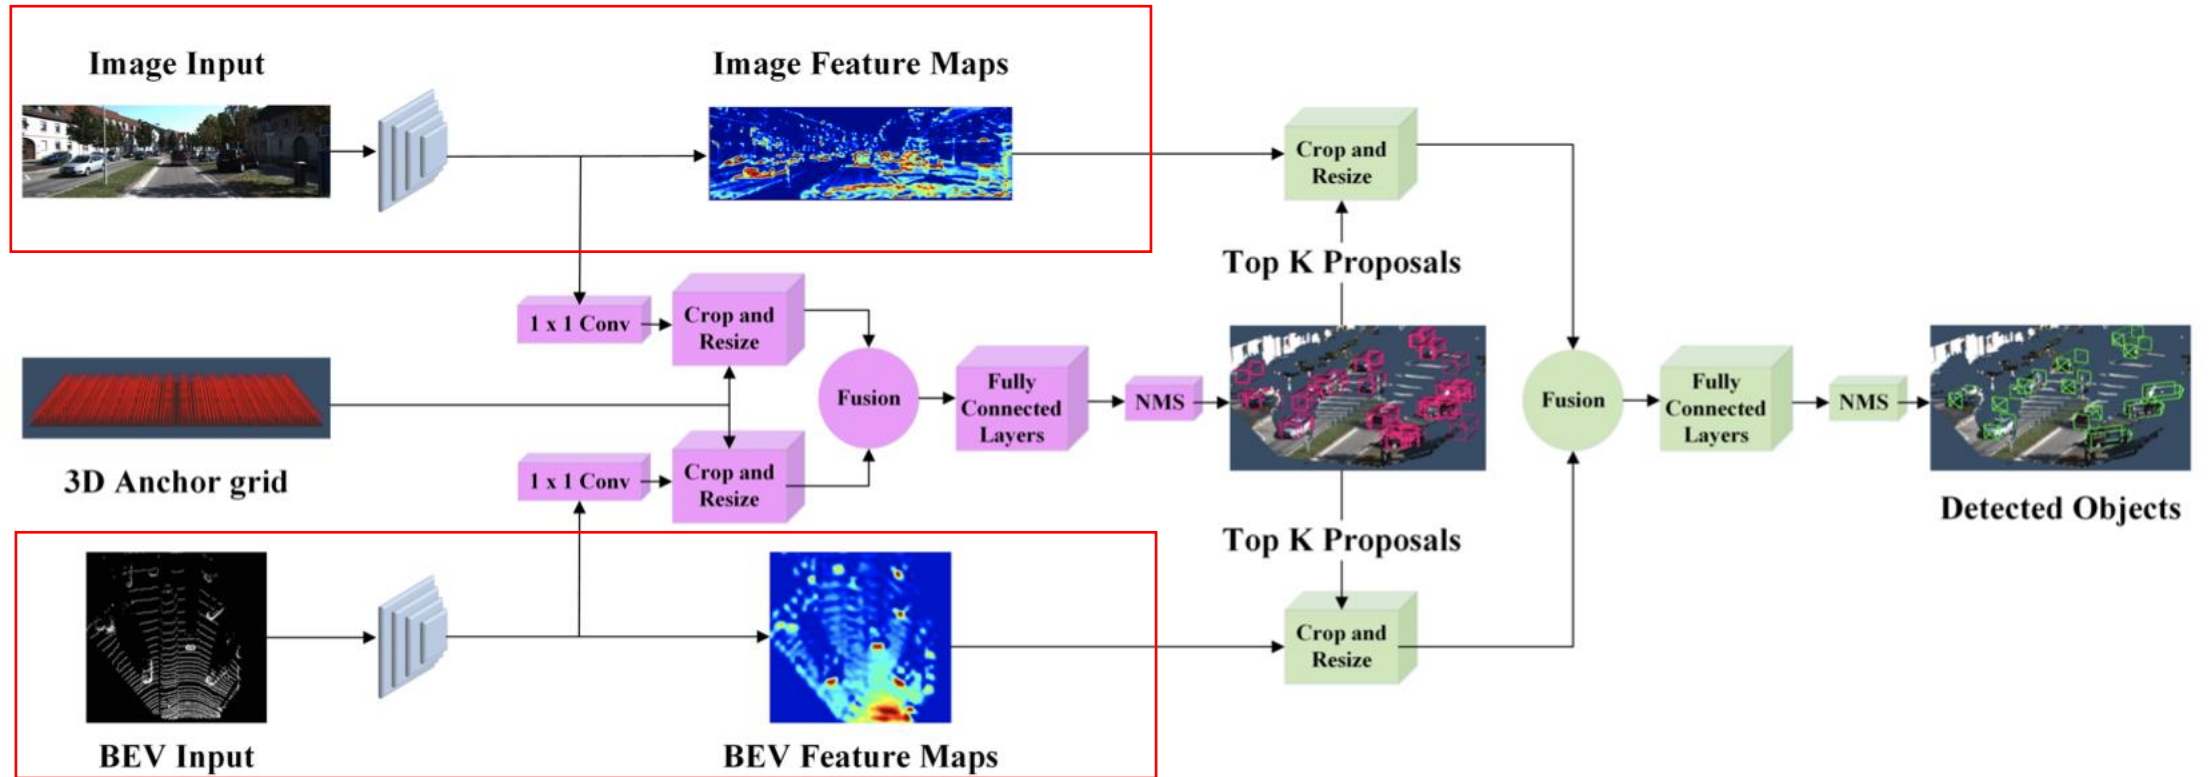



Figure 3.1: AVOD's architectural diagram. The feature extractors are shown in **blue**, the region proposal network in **pink**, and the second stage detection network in **green**.

Point Cloud Representation



reason:

- 3D grid representation requires complex and extensive computation for feature extraction

BEV maps:

- discretizing the point cloud with a 0.1 meters resolution
- projecting the voxels onto the xz-plane.
- the height feature is computed as the maximum height of the points in the cell
- point cloud is divided into 5 equally-sized slices

output 800*700*6:

- first 5 channels: maximum height of points in grid cell
- sixth channel: point density information,
normalized by $\min(1.0, \log(N + 1) / \log 16)$
——N is the number of points in the cell

Generating Feature Maps——Extractor

Operation	Kernel	Output	Operation	Kernel	Output
conv1	3×3	$480 \times 1590 \times 32$	conv1	3×3	$700 \times 800 \times 32$
maxpool	2×2	$240 \times 795 \times 32$	maxpool	2×2	$350 \times 400 \times 32$
conv2	3×3	$240 \times 795 \times 64$	conv2	3×3	$350 \times 400 \times 64$
maxpool	2×2	$120 \times 397 \times 64$	maxpool	2×2	$175 \times 200 \times 64$
conv3	3×3	$120 \times 397 \times 128$	conv3	3×3	$175 \times 200 \times 128$
maxpool	2×2	$60 \times 198 \times 128$	maxpool	2×2	$87 \times 100 \times 128$
conv4	3×3	$60 \times 198 \times 256$	conv4	3×3	$87 \times 100 \times 256$
upsampling	NA	$240 \times 795 \times 256$	upsampling	NA	$350 \times 400 \times 256$
1×1 conv	1×1	$240 \times 795 \times 1$	1×1 conv	1×1	$350 \times 400 \times 1$

VGG16:

- Half filters
- Xavier weight initialize
 - $1/N_i$
 - $1/N_{i+1}$
 - $2/(N_i + N_{i+1})$
- Batch normalization
- Discard the fourth maxpooling


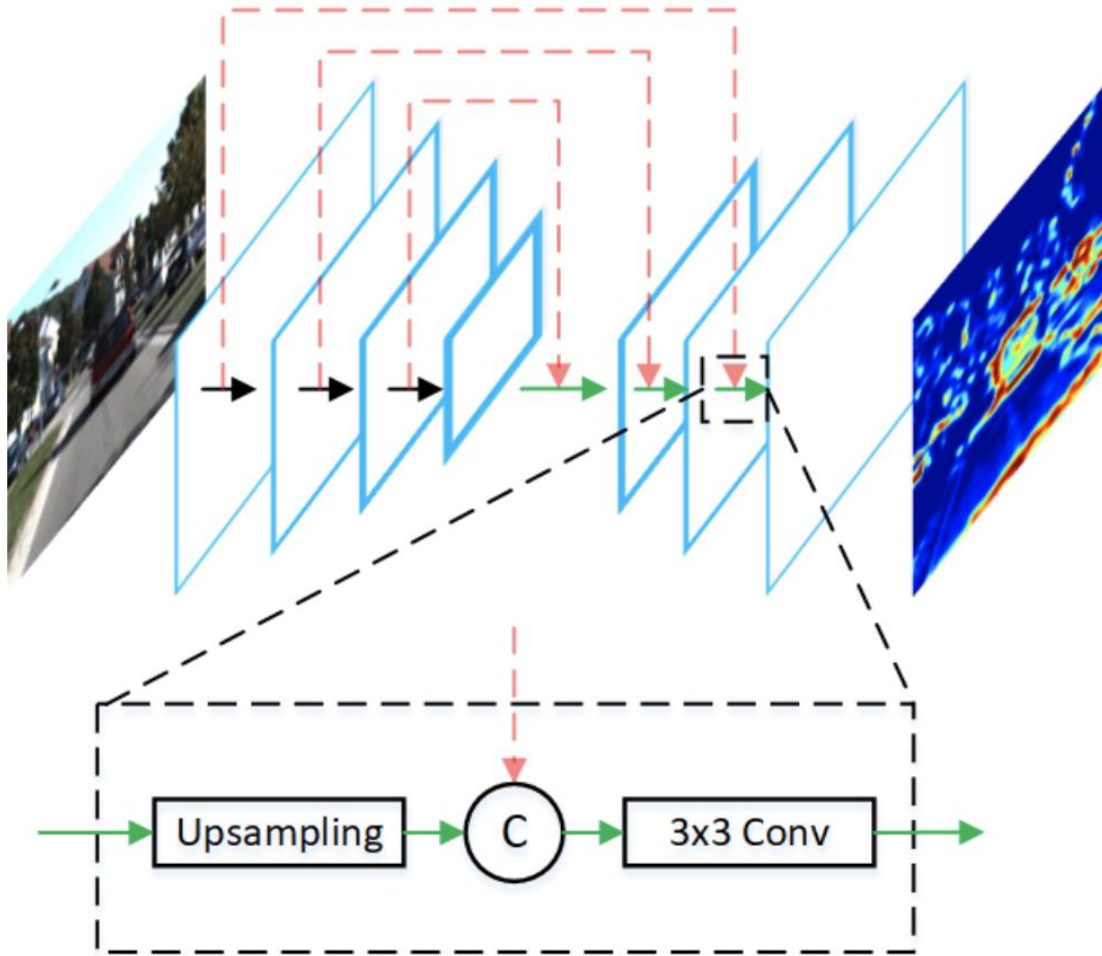
Upsampling

- 4* bilinear upsampling layer

Image branch with input of $(480 \times 1590 \times 3)$

BEV branch with input of $(700 \times 800 \times 6)$

Generating Feature Maps—Pyramids



advantage:

- The final feature map is of high resolution and representational power
- can significantly boost the performance of the network for detecting small objects

RPN Network

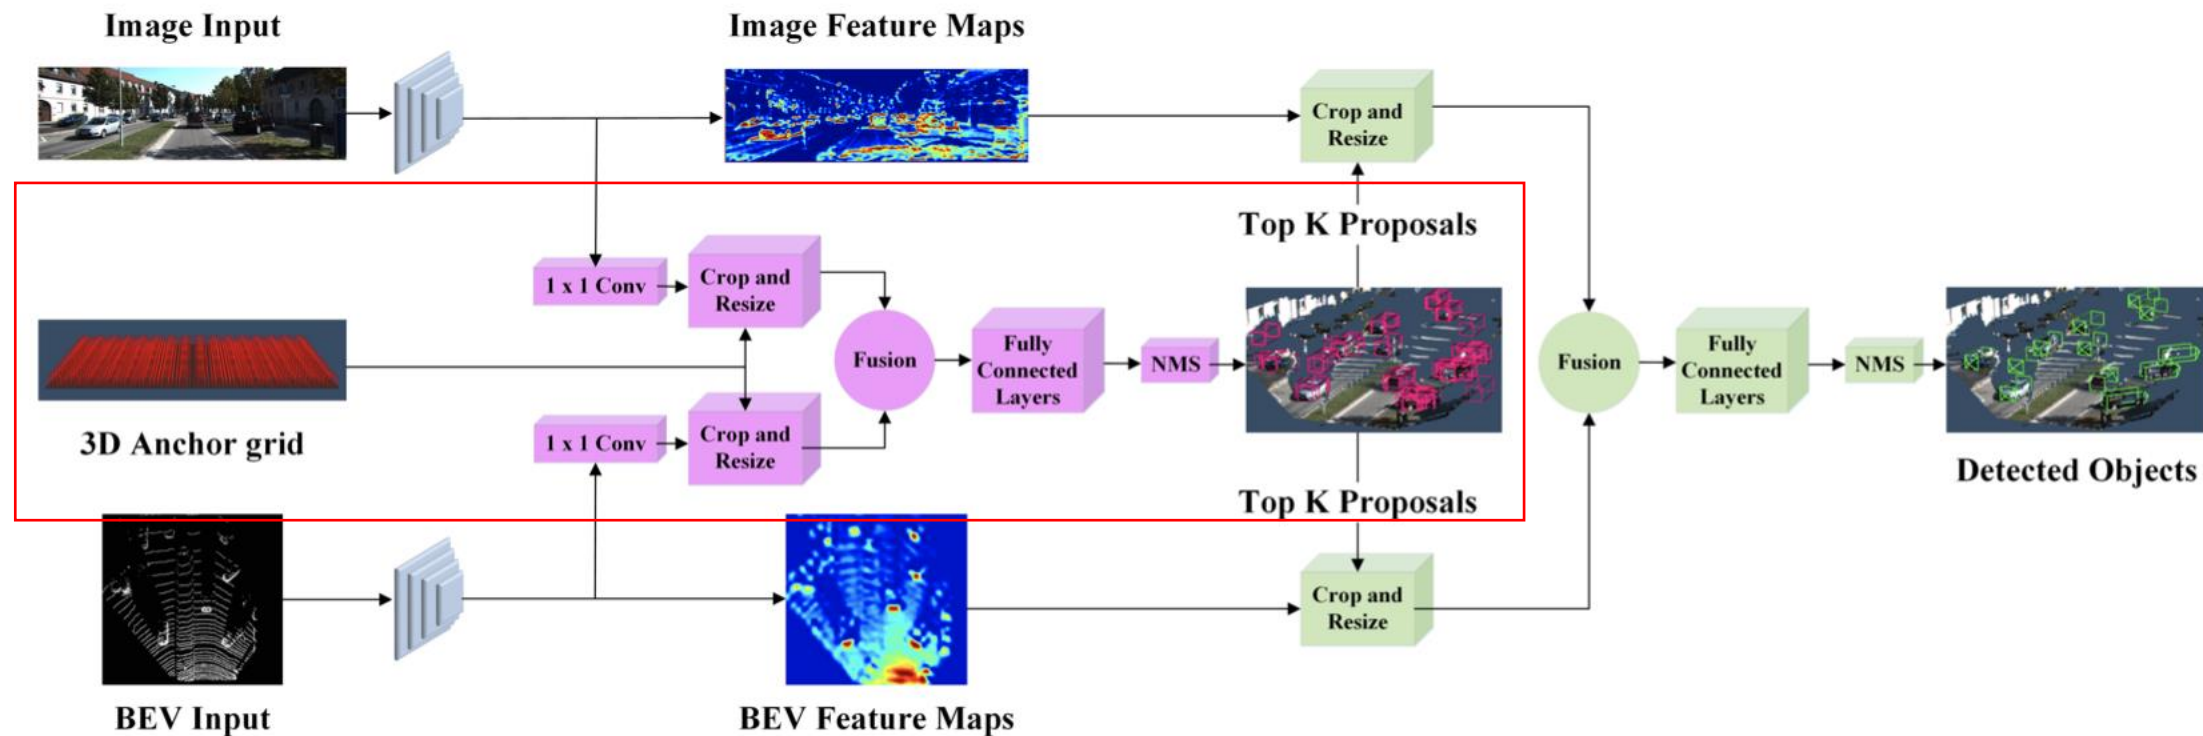





Figure 3.1: AVOD's architectural diagram. The feature extractors are shown in **blue**, the region proposal network in **pink**, and the second stage detection network in **green**.

Anchor Generation



Anchor Generation

parameter:

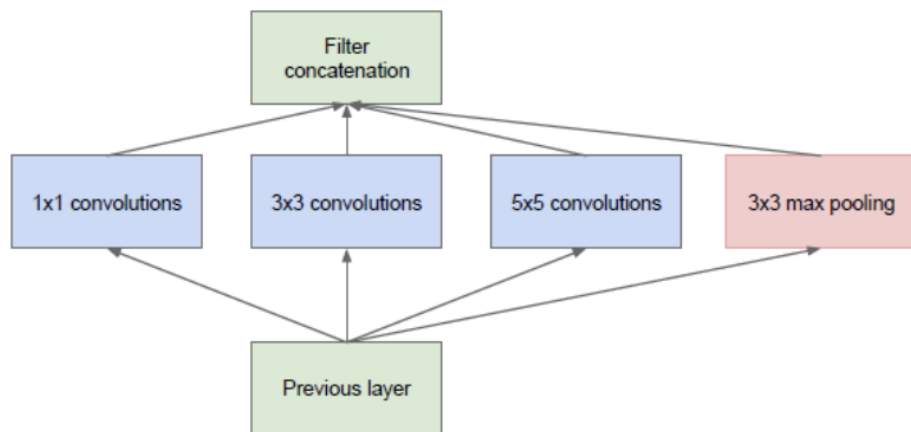
- centroid (tx, ty, tz) + dimensions (dx, dy, dz)
- (tx, ty) : sampled at an interval of 0.5 meters in BEV
- tz : based on the sensor's height
- (dx, dy, dz) : determined by clustering the training samples for each class
- **orientations**: generated with 0° and 90°

Reason for using clusters:

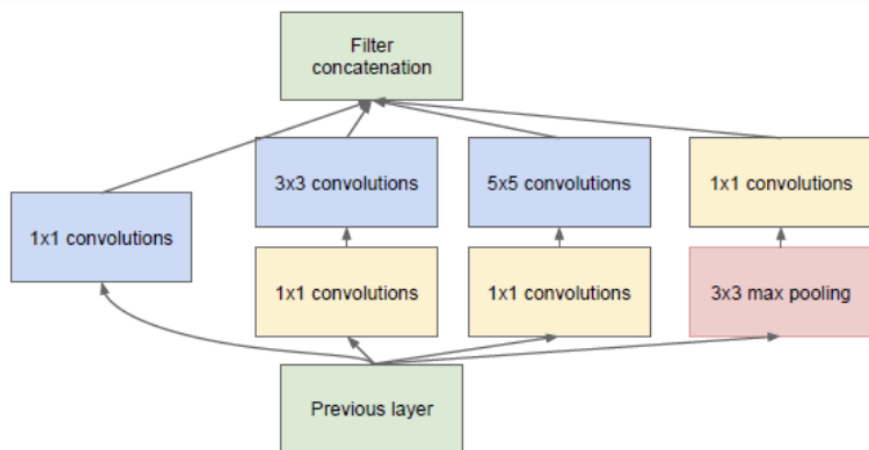
- hard to hand pick the best prior box dimensions

Result in 10~100K non-empty anchors per frame

Dimensionality Reduction Via 1*1 Convolutional Layers



(a) Inception module, naïve version



(b) Inception module with dimension reductions

作用:

- 实现跨通道的交互和信息整合
- 进行卷积核通道数的降维和升维
- 可以实现与全连接层等价的效果
- 减少内存，提高运算速度

GoogLeNet:


feature map= $28 \times 28 \times 192$

(a) $1 \times 1 \times 192 \times 64 + 3 \times 3 \times 192 \times 128 + 5 \times 5 \times 192 \times 32$

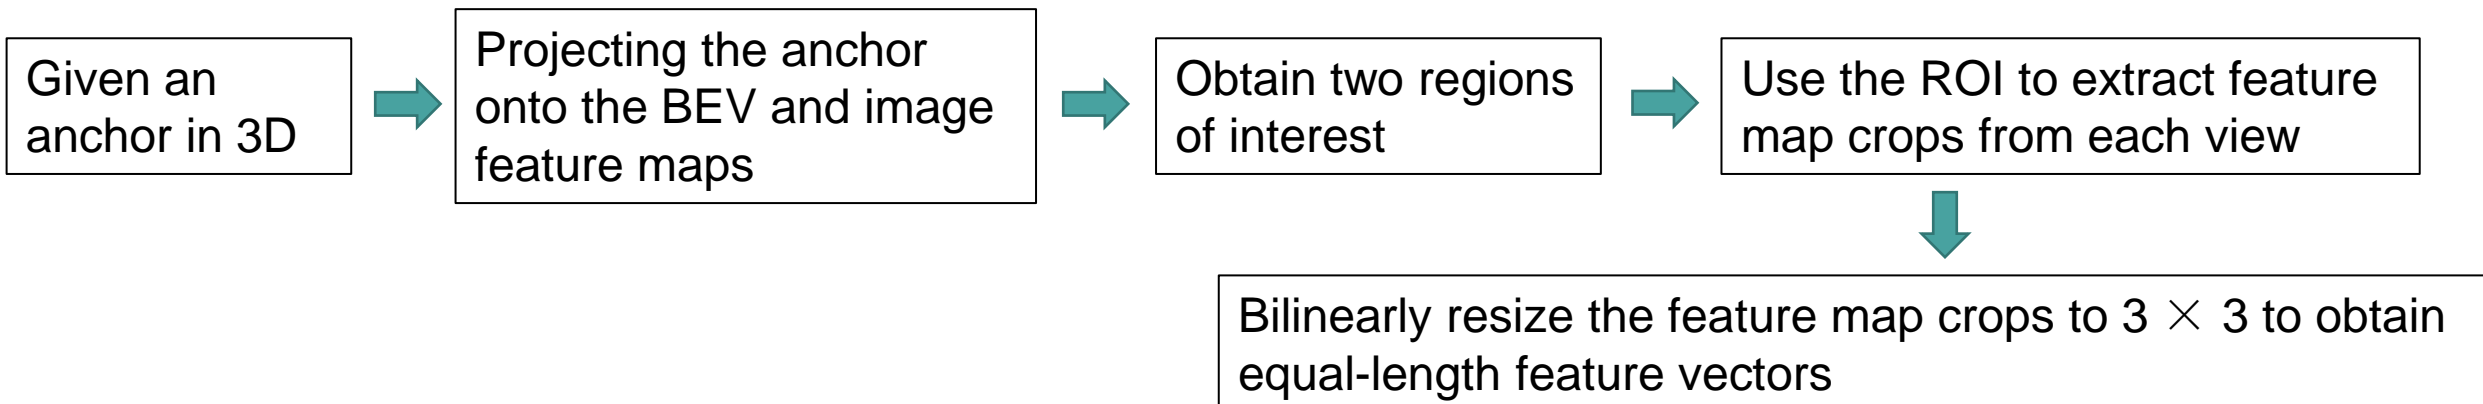
(b) $1 \times 1 \times 192 \times 64 + (1 \times 1 \times 192 \times 96 + 3 \times 3 \times 96 \times 128)$
 $+ (1 \times 1 \times 192 \times 16 + 5 \times 5 \times 16 \times 32)$

参数大约减少到原来的三分之一。


Extracting Feature Crops Via Multiview Crop And Resize Operations:



```
# Do ROI Pooling on BEV
bev_proposal_rois = tf.image.crop_and_resize(
    bev_proposal_input,
    self._bev_anchors_norm_pl,
    tf_box_indices,
    self._proposal_roi_crop_size)
# Do ROI Pooling on image
img_proposal_rois = tf.image.crop_and_resize(
    img_proposal_input,
    self._img_anchors_norm_pl,
    tf_box_indices,
    self._proposal_roi_crop_size)
```



3D Proposal Generation:



$$L(\{p_i\}, \{t_i\}) = \frac{1}{N_{cls}} \sum_i L_{cls}(p_i, p_i^*) + \lambda \frac{1}{N_{reg}} \sum_i p_i^* L_{reg}(t_i, t_i^*)$$

cross-entropy loss
for classification

Smooth L1 loss
for 3D box regression

background anchors
are ignored $\Rightarrow p_i^* = 0$

- **Fusion:** fuse the feature crops from both views via an **element-wise mean operation**.
- **Fully Connected Layers: size 256**
 - **regress axis aligned object proposal boxes:** computing $(\Delta t_x, \Delta t_y, \Delta t_z, \Delta d_x, \Delta d_y, \Delta d_z)$, the difference in centroid and dimensions between anchors and ground truth bounding boxes.
 - **output an object/background “objectness” score:** calculating the **2D IoU in BEV** between the anchors and the ground truth bounding boxes.

object/background
anchors

For the car class: less than **0.3** \Rightarrow background anchors(negative)
greater than **0.5** \Rightarrow object anchors(positive)

For the pedestrian and cyclist classes: greater than **0.45**
 \Rightarrow object anchors (positive)

- **NMS:** To remove redundant proposals; IoU threshold of **0.8** in BEV
- **Output:** 300 proposal for car class, 1024 proposals for pedestrian and cyclist.

Second Stage Detection Network

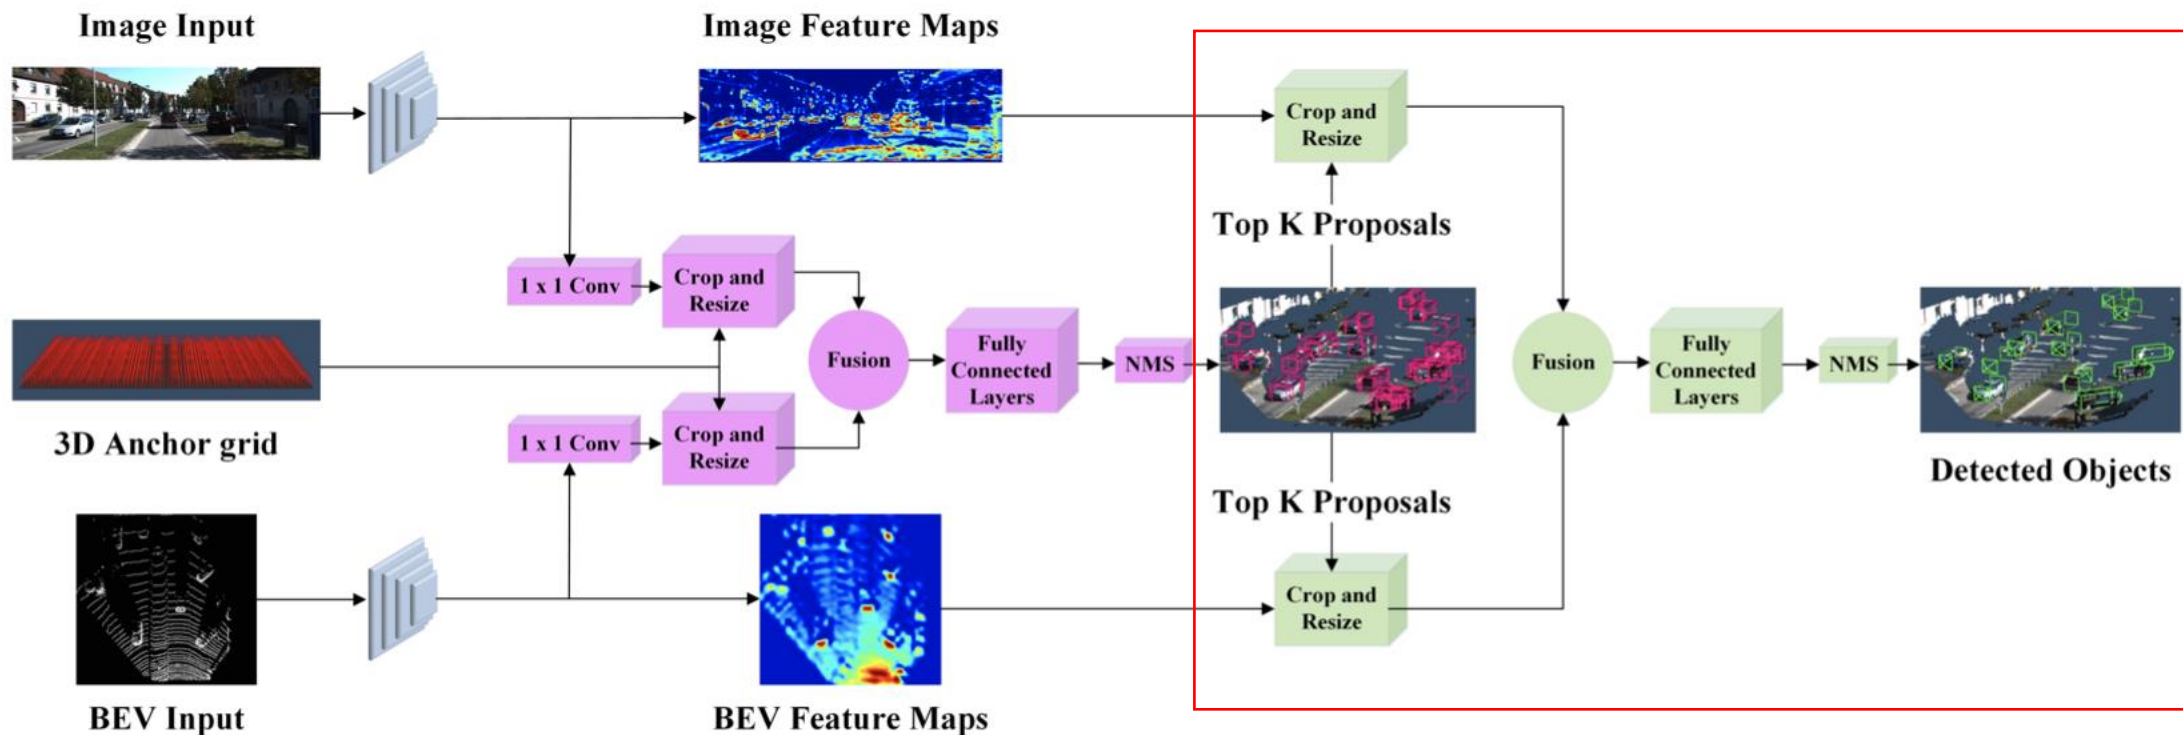


Figure 3.1: AVOD's architectural diagram. The feature extractors are shown in **blue**, the region proposal network in **pink**, and the second stage detection network in **green**.

3D Bounding Box Encoding




Fig. 4: A visual comparison between the 8 corner box encoding proposed in [4], the axis aligned box encoding proposed in [16], and our 4 corner encoding.

Reduces the box representation from an over parameterized **24** dimensional vector to a **10** dimensional one.

Regression targets:

$$(\underbrace{\Delta x_1 \dots \Delta x_4, \Delta \bar{y}_1 \dots \Delta \bar{y}_4}_{\text{correspond to the closest corner of the proposals to the closest corner of the ground truth box}}, \underbrace{\Delta \bar{h}_1, \Delta h_2}_{\text{from the ground plane between the proposals and the ground truth boxes}})$$

correspond to the closest corner of the proposals to the closest corner of the ground truth box

from the ground plane between the proposals and the ground truth boxes

Explicit Orientation Vector Regression


MV3D relies on the extents of the estimated bounding box where the orientation vector is assumed to be in the **direction of the longer side of the box**

Two problems:



1. fails for detected objects like pedestrians:
2. Orientation information is lost as the corner order is not preserved in closest corner to corner matching: $\pm\pi$ radians

Solution:
Computing $(x_{or}, y_{or}) = (\cos(\theta), \sin(\theta))$

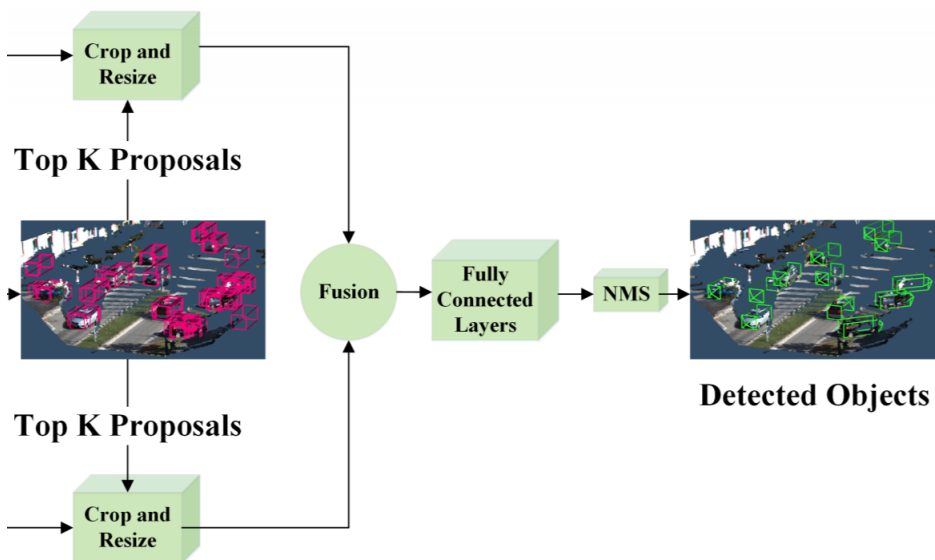
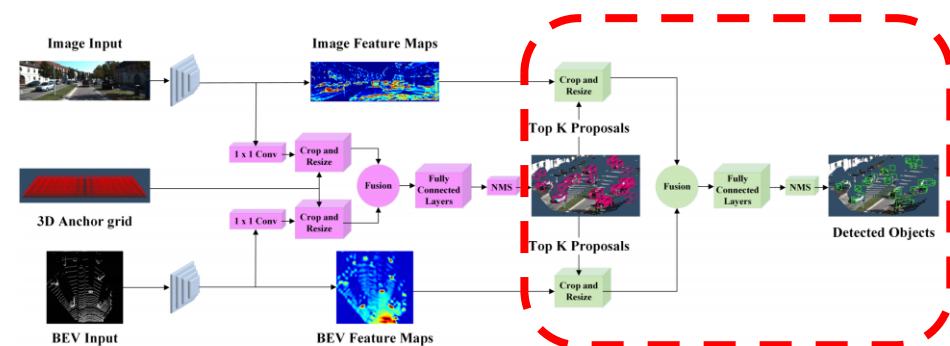
Every $\theta \in [-\pi, \pi]$ can be represented by a unique unit vector in the BEV space.



- the regressed orientation vector
- the possible headings
the closest heading
- the possible headings



Generating Final Detections



Input: feature crops generated from projecting the **Top K** Proposals into the two input views.

Process:

- 1、 **Crop and Resize:** resize the crops to 7x7
- 2、 **Fusion:** fused with an element-wise **mean** operation
- 3、 **Fully Connected Layers:** three fully connected layers of size 2048
- 4、 **NMS:** To remove overlapping detections, set a threshold of 0.01.

Loss Function:

$$L(\{p_i\}, \{t_i\}) = \frac{1}{N_{cls}} \sum_i L_{cls}(p_i, p_i^*) + \lambda \frac{1}{N_{reg}} \sum_i p_i^* L_{reg}(t_i, t_i^*) + \lambda \frac{1}{N_{ang}} \sum_i p_i^* L_{ang}(t_i, t_i^*)$$

two Smooth L1 losses for the bounding box and orientation vector regression tasks, and a cross-entropy loss for the classification task.

Notes:

Proposals 2D IoU in BEV with the ground truth boxes

For **car**: at least **0.65**

For **pedestrian/cyclist**: at least **0.55**

Result

Method	Runtime (s)	Class	AP_{3D} (%)			AP_{BEV} (%)		
			Easy	Moderate	Hard	Easy	Moderate	Hard
MV3D [4]	0.36	Car	71.09	62.35	55.12	86.02	76.90	68.49
VoxelNet [10]	0.23		77.47	65.11	57.73	89.35	79.26	77.39
F-PointNet [12]	0.17		81.20	70.39	62.19	88.70	84.00	75.33
Ours	0.08		73.59	65.78	58.38	86.80	85.44	77.73
Ours (Feature Pyramid)	0.1		81.94	71.88	66.38	88.53	83.79	77.90
VoxelNet [10]	0.23	Ped.	39.48	33.69	31.51	46.13	40.74	38.11
F-PointNet [12]	0.17		51.21	44.89	40.23	58.09	50.22	47.20
Ours	0.08		38.28	31.51	26.98	42.51	35.24	33.97
Ours (Feature Pyramid)	0.1		46.35	39.00	36.58	50.66	44.75	40.83
VoxelNet [10]	0.23	Cyc.	61.22	48.36	44.37	66.70	54.76	50.55
F-PointNet [12]	0.17		71.96	56.77	50.39	75.38	61.96	54.68
Ours	0.08		60.11	44.90	38.80	63.66	47.74	46.55
Ours (Feature Pyramid)	0.1		59.97	46.12	42.36	62.39	52.02	47.87



PART 3

AVOD code

<https://github.com/kujason/avod>

Dataset and code architecture


```
Kitti
  object
    testing
    training
      calib
      image_2
      label_2
      planes
      velodyne
    train.txt
    val.txt
```

- Generate mini-batches for the RPN
- Train on the specific config
- Run Evaluator
- Run inference on the val split
- Viewing Results

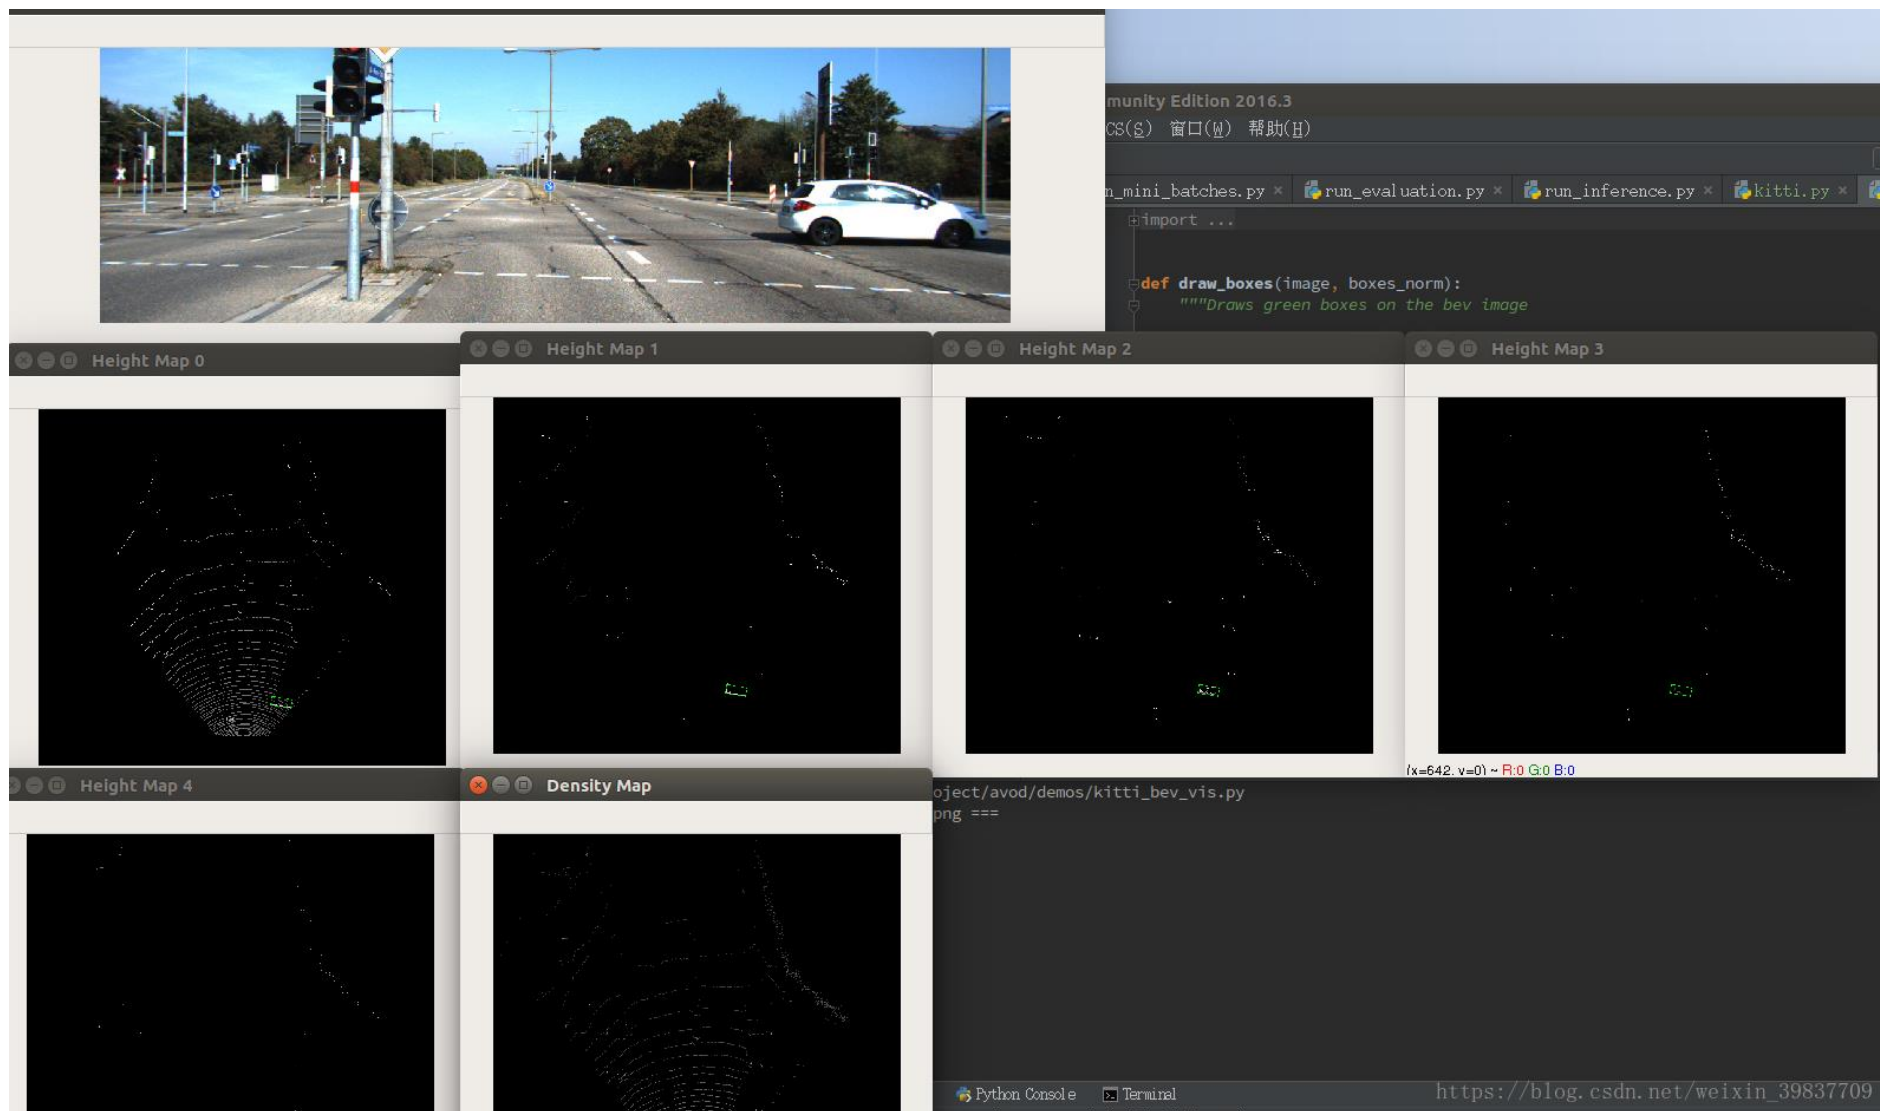
```
avod ~/Desktop/store/avod
└─ avod
  ├── builders
  ├── configs
  ├── core
  ├── data
  ├── datasets
  ├── experiments
  ├── protos
  ├── tests
  ├── utils
  ├── __init__.py
  └─ demos
    ├── augmentation
    ├── dataset
    ├── generate_anchors.py
    ├── kitti_bev_vis.py
    └── show_predictions_2d.py
```

```
scripts
├─ install
├─ offline_eval
├─ preprocessing
├─ __init__.py
└─ wavedata
  ├── demos
  ├── scripts
  └─ wavedata
    ├── .coveragerc
    ├── .gitignore
    ├── .travis.yml
    ├── LICENSE
    ├── README.md
    └── requirements.txt
.coveragerc
.gitignore
.gitmodules
.travis.yml
LICENSE
README.md
requirements.txt
```


kitti.py 运行效果



kitti_bev_vis.py 运行效果



—END—
THANK YOU

